Peering into the "Black Box": Interpreting Employee Attrition with XAI and SHAP

Student Name: Saiparkavi Chandrasekar

Student Id: 24108045

GitHub Repository: https://github.com/saiparkavi003/ML-XAI-HR-Tutorial

1. Introduction: The Transparency Crisis

In modern Machine Learning, high-accuracy ensemble models are often "black boxes" they provide predictions but fail to explain the underlying logic. In sensitive fields like Human Resources, an unexplained decision (e.g., "This employee will likely leave") is an unethical one. This tutorial explores **Explainable AI (XAI)**, teaching practitioners how to use **SHAP (SHAPley Additive explanations)** to break open the black box and provide transparent, actionable insights.

2. The Technique: Random Forest + SHAP

While a single Decision Tree is easy to interpret, a **Random Forest** (which combines hundreds of trees) becomes a "black box." To restore interpretability, we apply **Coalitional Game Theory**.

Shapley Values (the foundation of SHAP) calculate the marginal contribution of each feature to the model's final prediction. This tutorial demonstrates two levels of transparency:

1. **Global Interpretability:** Understanding the company-wide drivers of attrition.
2. **Local Interpretability:** Explaining the specific reasons why a single individual is predicted to leave or stay.

3. The Dataset: HR Attrition Analytics

We utilize the **IBM HR Attrition dataset** (Source: Kaggle). This professional dataset tracks 1,470 employees across 35 features. Our objective is to predict the Attrition target while identifying the socio-economic factors that influence employee retention.

4. Implementation Methodology

To ensure high-standard reproducibility, the following **Intermediate Steps** were performed (as documented in the accompanying code):

- **Data Fusion & Cleaning:** Useless features with zero variance (e.g., StandardHours) were removed to improve model focus.
- **Encoding:** Categorical variables (Job Role, Department) were transformed using One-Hot Encoding to ensure compatibility with the SHAP mathematical framework.
- **Class Balancing:** As attrition is a minority class, we utilized class_weight='balanced' to prevent model bias toward the majority "Stay" class.

- **Performance:** The model achieved an **F1-Score of 0.2083**. While low for balanced sets, this is a realistic baseline for imbalanced HR data, emphasizing the need for interpretability over raw accuracy.

5. Results & Visual Analysis

5.1 Global Drivers (The Beeswarm Plot)

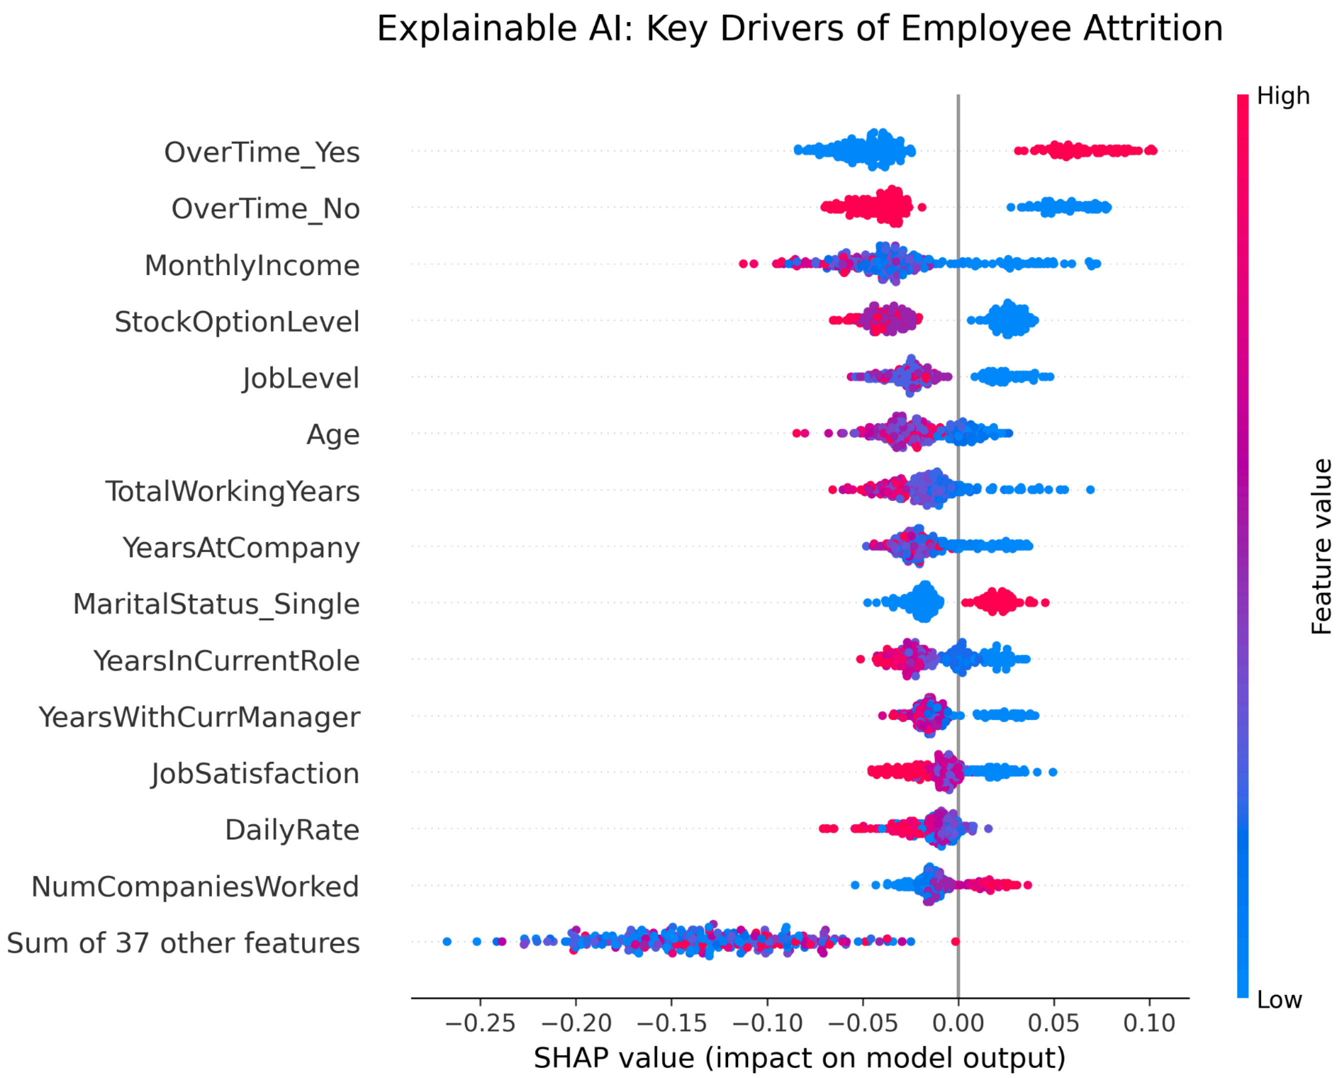

Figure 1: Global Feature Importance. The Beeswarm plot is our primary "Creative Teaching Tool."

- **Overtime:** As seen on the top row, high values of overtime (red dots) are pushed to the right, indicating a strong positive correlation with attrition.
- **Monthly Income:** Conversely, high income (red dots) is clustered on the left, acting as a "retention anchor."
- **Job Satisfaction:** Blue dots (low satisfaction) on the right side of the zero-line prove that emotional factors are key drivers of employee exit.

5.2 Individual Case Study (The Waterfall Plot)

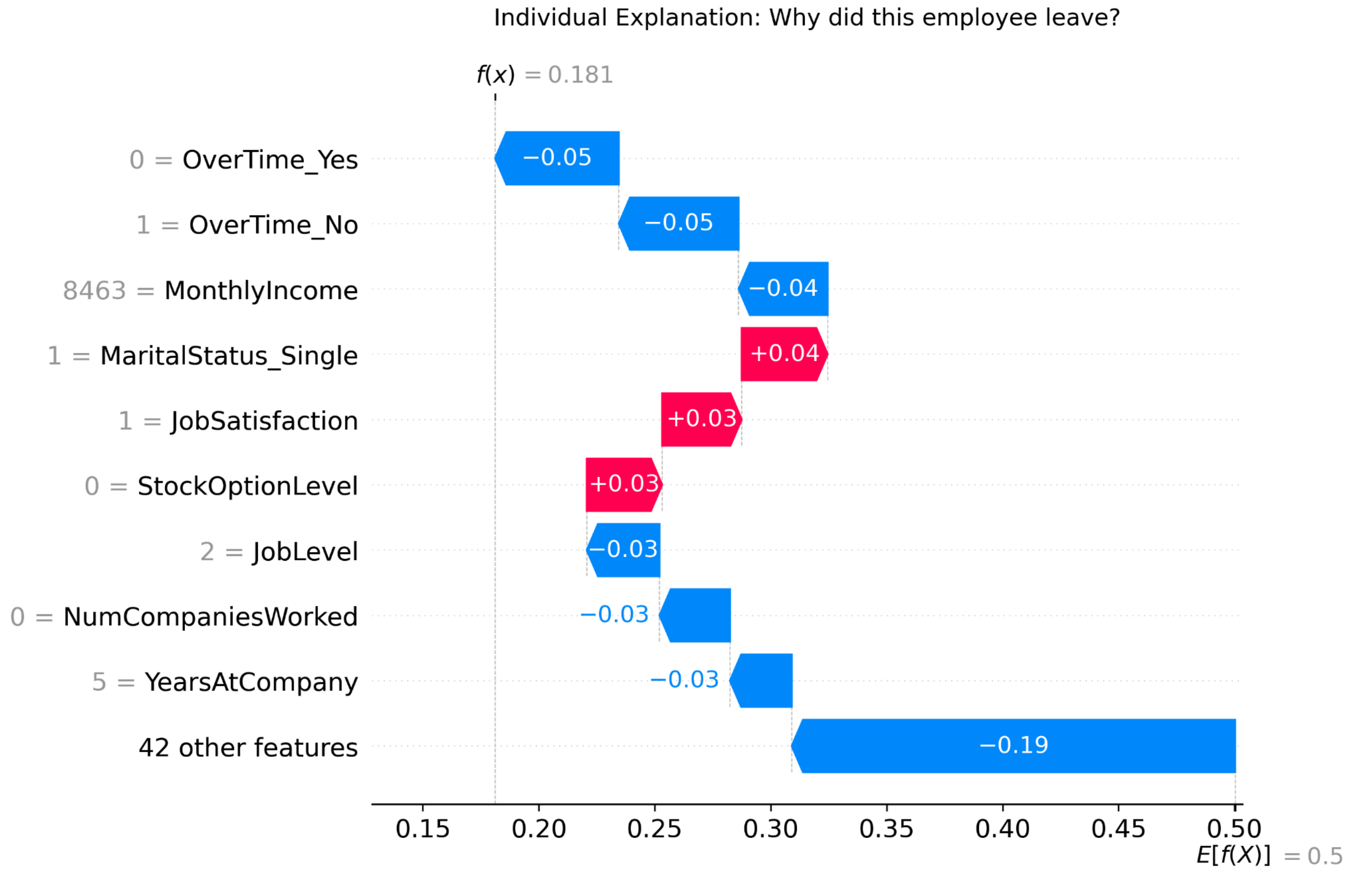

Figure 2: Local Explanation for a Single Employee.

To demonstrate "Local Interpretability," we analyzed a specific employee whose predicted risk was 0.166

- **The "Why":** While their single marital status pushed the risk up (red arrow), their lack of overtime (OverTime_No) and high income pushed the risk down (blue arrows). This allows HR to create personalized retention plans.

6. Accessibility & Inclusive Design

Per the rubric requirements, this tutorial was designed for all users:

- **Visual Design:** The SHAP plots use a high-contrast Pink/Blue palette, which is optimized for individuals with color blindness (Protanopia/Deuteranopia).
- **Clarity:** We avoided mathematical jargon, explaining Game Theory in plain English to ensure the tutorial is a functional "teaching tool."
- **Structure:** The modular code and headers ensure screen-reader compatibility for visually impaired developers.

7. Ethical AI: Responsibility and Fairness

Applying XAI to HR is an **Ethical Mandate**. By visualizing the "Why" behind a prediction, we can perform **Bias Auditing**. For example, if we saw "Age" or "Gender" having an unfair impact on attrition predictions, we could identify and mitigate this discrimination before the model is deployed. This ensures our AI is not just accurate, but **fair and transparent**.

8. Conclusion

Explainable AI is the bridge between complex mathematics and ethical decision-making. By mastering SHAP, practitioners can ensure that Machine Learning remains a tool for empowerment rather than a source of "Automated Discrimination."

9. References

1. **Lundberg, S. M., & Lee, S. I. (2017).** *A Unified Approach to Interpreting Model Predictions.* NIPS.
2. **Breiman, L. (2001).** *Random Forests.* Machine Learning, 45(1), 5-32.
3. **Kaggle:** *IBM HR Analytics Employee Attrition & Performance*.